

**Deep Learning (CSL7590)**  
**Assignment 2**



**Submitted By:**

Shivani Tiwari (M24CSA029)  
Suvigya Sharma (M24CSA033)

**Submitted To:**

Dr. Deepak Mishra

**Collab link:**

[https://colab.research.google.com/drive/1PPocx\\_-sBkQ9U0x4MZmveJ\\_XPs-v3pR9?usp=sharing](https://colab.research.google.com/drive/1PPocx_-sBkQ9U0x4MZmveJ_XPs-v3pR9?usp=sharing)

**Objective:** The goal of this assignment is to build a single Convolutional Neural Network (CNN) in PyTorch that can classify images from the CIFAR-100 dataset into three different categories at the same time:

1. Fine-level classification – Identify the exact object from 100 possible classes.
2. Superclass classification – Group the object into one of 20 broader categories.
3. Synthesized group classification – Further categorize the object into one of 9 custom-defined groups based on shared characteristics (e.g., vehicles, aquatic animals, plants, etc.).
  - Plants/Parts of Plants → Superclasses: flowers, trees, fruits and vegetables
  - Vehicles → Superclasses: vehicles1, vehicles2
  - Invertebrates → Superclasses: non-insect invertebrates, insects
  - Aquatic Animals → Superclasses: fish, aquatic mammals
  - Large Animals → Superclasses: large carnivores, large omnivores and herbivores
  - Man-made Articles → Superclasses: food containers, household electrical devices, household furniture, large man-made outdoor things
  - People → Superclass: people
  - Normal Terrestrial Animals → Superclasses: reptiles, medium-sized mammals, small mammals
  - Outdoor Scenes → Superclass: large natural outdoor scenes
4. To improve performance and handle class imbalances, we introduce a severity-based loss function that penalizes misclassifications based on their severity. The model is trained and evaluated across multiple train-test splits (70:30, 80:20, and 90:10).

## **Methodology:**

### **1. Data Preprocessing:**

- The CIFAR-100 dataset is loaded and preprocessed using normalization.
- A custom PyTorch dataset class (CIFAR100\_Custom) is created to map each fine class to its respective superclass and group category.
- The dataset is split into training and testing sets with three different ratios (70:30, 80:20, 90:10).
- To handle data imbalance, we use a weighted sampler, ensuring that underrepresented classes contribute equally to training.

### **2. CNN Architecture model:**

- Feature Extraction Layers: 3 sets of convolutional layers with Batch Normalization and ReLU activation. Max-pooling layers to reduce dimensionality.
- Fully Connected Layers: Two dense layers with Batch Normalization and Dropout to prevent overfitting.
- Three Classification Heads: One output layer for each classification task (fine class, superclass, and group classification).

### **3. Loss Function & Training Process**

The cross-entropy loss is used for classification.

A severity penalty matrix is introduced to assign different levels of penalty based on the type of misclassification:

- Same superclass → Minor penalty.
- Same group but different superclass → Moderate penalty.
- Different group → Highest penalty.

The final severity-weighted loss function adjusts training based on these penalties.

### **Training Process:**

- The model is trained for 20 epochs using the Adam optimizer (learning rate = 0.001).
- The dataset is divided into three train-test splits to analyze model performance across different training size.

#### 4. Evaluation & Performance Metrics

- After training, the model is tested on the validation dataset.
- Accuracy is computed for all three classification heads (fine class, superclass, and group classification).
- Loss trends over training epochs
- Confusion matrices are plotted for each classification task to visualize misclassifications.
- The total number of trainable and non-trainable parameters.

#### Results

Total Trainable Parameters: 13488001

Split Ratio	Final Loss	Class Accuracy	Superclass Accuracy	Group Accuracy
70:30	0.492	52.06	52.55	54.95
80:20	0.4829	53.56	53.79	57.44
90:10	0.459	55.65	56.01	59.25

#### Observations

- Accuracy improves with larger training sets: As the training data increases, model accuracy improves across all classification tasks.
- Group classification performs better than fine classification: This suggests that distinguishing broad categories is easier than differentiating fine-grained classes.
- Loss reduction is significant over epochs: The final loss values show significant improvement over training.

### Bonus part: Enhancing CNN Training with Severity-Based Misclassification Penalties

Split Ratio	Final Loss	Class Accuracy	Superclass Accuracy	Group Accuracy
70:30	0.3622	55.31	55.18	58.04
80:20	0.3507	56.75	56.67	59.05
90:10	0.4401	59.52	59.36	62.48

### Observations

1. Bonus approach reduces severe misclassification: Accuracy gains indicate a lower frequency of misclassification across groups.
2. Computational cost increases slightly: Since the loss function includes additional penalty terms, training takes marginally longer.
3. Fine-class accuracy benefits the most: The model learns to make more precise predictions within the superclass and groups

## References

1. <https://www.kaggle.com/datasets/fedesoriano/cifar100>
2. <https://www.geeksforgeeks.org/introduction-convolution-neural-network/>
3. <https://www.coursera.org/learn/convolutional-neural-networks>
4. <https://pytorch.org/docs/stable/nn.html>
5. <https://www.geeksforgeeks.org/handling-imbalanced-data-for-classification/>